

Phase A: Ethical Analysis

Grounded in Task 6: Constructing and Evaluating Synthetic Media

1. Return to What I Built

Going back to the two artifacts with a few weeks of distance, the thing that strikes me most isn't the video or the audio itself - it's how little friction there was in producing something that could pass as a real recorded human. The ElevenLabs narration (artifact 1) took about five minutes to generate from an unmodified script, required no special skill, and turned out convincing enough that a purpose-built detector - deepfakedetection.io, scored it 95% "Likely Authentic." A tool designed for exactly one job, telling synthetic speech apart from real speech, got it backwards, and did so with a confident, detailed-sounding explanation that on closer inspection was partly generic template language rather than a real per-file analysis.

That result is unsettling for a specific reason: the content of the narration was true - it was a fact-checked summary of my own Task 5 analysis, and the disclosure was present throughout the repo. Nothing about the artifact was designed to deceive anyone. And yet the mechanism that produced it - a script, a stock voice, a few minutes of processing is identical to the mechanism that would produce a fabricated statement from a real, named person. Truthfulness and disclosure lived entirely outside the technical pipeline; the pipeline itself was indifferent to both.

A few things the process log doesn't fully capture:

- **The gap between "impressive" and "trustworthy."** Listening to the ElevenLabs audio again, I notice I have no vocabulary for the thing that made it feel slightly off in a few stretches - no mispronunciation, no flat cadence, nothing nameable. It's a feeling that only exists because I already know it's synthetic. A listener without that prior knowledge has no foothold for it at all, which is exactly the finding that should be alarming rather than reassuring.
- **Where the tools "refused" - or didn't.** Nothing in this task triggered a content warning or refusal from any tool. ElevenLabs, LOVO, and D-ID all generated exactly what was asked, using stock voices and a stock avatar, without prompting for any confirmation that the content was truthful or that I had rights to say it in this way. The closest thing to friction I hit was commercial, not ethical: TTSMaker and Murf.ai both let me preview or generate audio for free but gated the actual file download behind a paid tier.
- **What surprised me.** Not the video - the video's flaws (excess avatar motion, a slightly mismatched voice/avatar pairing, flat emotional affect) were visible enough that I'd expect a casual viewer to register something as "off" within a few seconds, even without being able to say why. The audio is what surprised me, because it had no equivalent tell. If I'd stopped after Artifact 1 and never run it against a detector, I would have had no independent way to know it wasn't real.
- **Would I build this again?** Yes, as constructed - a disclosed, consented, low-stakes artifact built to learn what the tools can and can't do. What I wouldn't do is treat the detection result casually. Finding out that a dedicated detection tool can be this wrong is the part of this task I take most seriously going into Phase B, because it means "run it through a detector" is not a mitigation I can recommend anyone rely on.

2. Reasoning Across the Axes

Truth axis

Same delivery mechanism, false content.

Scenario: A student government candidate at a university loses a close election. A week later, an anonymous account posts an “audio recording” of the winning candidate - generated the same way I generated Artifact 1, same category of consumer voice tool, same absence of any technical barrier - in which the candidate appears to admit that a campaign promise was a lie told deliberately to win votes. The clip is eleven seconds long, sounds exactly as natural as my ElevenLabs narration did, and starts circulating in group chats before anyone official can respond.

What changes: Nothing about the *mechanism* changes at all - this is the same pipeline, possibly the same free tier of the same tool. What changes is that the content is no longer answerable to reality. My narration was true because I had verified it in Task 5; this clip is false because someone typed a lie into a text box. The technology has no way of knowing or caring which one it's asked to produce, and based on what I saw with TTSMaker, Murf.ai, ElevenLabs, LOVO, and D-ID - none of the tools I used asked. The ethical weight shifts entirely onto the person operating the tool and, once the clip escapes their control, onto everyone who shares it without asking whether an eleven-second recording is even the kind of thing that can be verified quickly. The most damage isn't done by people who believe the deepfake is real forever; it's done in the hours before anyone can debunk it, when the false content is doing its work uncontested - exactly the same window in which my true content, or a false one, would be equally convincing.

Consent axis

Same delivery, someone else's likeness, without consent.

Scenario: A small business owner has a public LinkedIn video where they discuss their company's values. Someone extracts thirty seconds of that speaker's voice, feeds it into a cloning tool with roughly the sophistication of ElevenLabs, and generates a new “recording” of that person seemingly authorizing a wire transfer to a new vendor account - sent as a voicemail to their own bookkeeper. This is structurally the same scam-call-to-a-grandparent scenario the assignment names, just moved into a business context I find easier to picture concretely.

What changes: This is the point where the technology stops being a tool and becomes a weapon, and the boundary is precisely consent. In my Task 6 work, every voice and avatar was either a licensed stock voice/avatar (Roger, Jenny, Amber) or my own verified content - nobody's actual identity was borrowed without their say. The bookkeeper scenario borrows a real, identifiable person's actual voice, from public material they posted for an entirely different purpose, and uses the trust that attaches to *that specific voice* to authorize an action the real person never took. The harm isn't abstract reputational damage - it's a direct request for money, backed by the same feeling of certainty (“that's clearly him, I recognize his voice”) that made my ElevenLabs narration pass a real detector at 95%. Once a real person's identity is doing the persuasive work without their knowledge, disclosure and truthfulness stop being sufficient safeguards, because the entire scam depends on the victim never being told it's synthetic.

Context axis

The disclosure label is stripped or the context is reframed.

Scenario: Someone downloads Artifact 2 (my D-ID coaching video) directly from the GitHub repo, trims out the parts of the narration that name Task 5 or reference the dataset, re-encodes it as a vertical video for a social platform, and reposts it captioned as genuine advice from “a lacrosse analytics expert” - no link back to my repo, no disclosure, no way for a viewer three shares downstream to discover the AI Voice/AI Video disclosure at the top of my original README.

What changes: My disclosure lives in the metadata layer around the artifact - a README banner, a file-naming convention (`_AI_VOICE_elevenlabs`, `_AI_VIDEO_did_lovo`) - not embedded irreversibly in the pixels or the waveform. Re-encoding for a different platform is a completely ordinary, non-malicious action (it happens to real, human-made content constantly), and it is sufficient, with no additional effort, to strip every trace of disclosure I built in. This is the axis where I have the least confidence any current mitigation actually holds: watermarks that survive re-encoding exist in principle (C2PA-style content credentials), but nothing in my process log involved verifying that any of ElevenLabs, LOVO, or D-ID actually embed tamper-resistant provenance data in their exports by default and if they don't, disclosure is only as durable as the first person who wants it gone.

Scale axis

One artifact over a task versus unlimited artifacts on demand.

Scenario: Instead of one script and two tools, imagine an actor with a list of two hundred candidates in local elections nationwide, a script template with the candidate's name and district swapped in programmatically, and API access to the same class of tools (ElevenLabs and D-ID both offer paid API tiers, not just the free consumer UI I used). In an afternoon, that actor could produce two hundred versions of the student-government scenario above, tailored to local issues in each race, released the week before two hundred separate elections.

What changes: My task took roughly 1.5–2 hours, most of it dead ends (failed TTSMaker and Murf.ai attempts, documentation) rather than actual generation time - the ElevenLabs step alone took about five minutes. At API scale, the marginal cost of the n -th fake statement approaches zero, and the review capacity of journalists, election officials, and platforms does not scale with it. A single well-made fake is a needle in a haystack that debunking can find; two hundred simultaneous fakes, each locally specific enough to be individually plausible, overwhelm any manual verification process that assumes it's dealing with rare events. Scale doesn't just mean “more of the same harm” - it changes which mitigations are even theoretically viable. Human fact-checking assumes scarcity; the tools I used this task assume abundance.

3. Mitigation Landscape Survey

Disclosure norms (labeling, watermarks, spoken acknowledgments). My own repo shows both what this gets right and where it breaks: an explicit README banner and a strict file-naming convention (`_AI_VOICE_elevenlabs`, `_AI_VIDEO_did_lovo`) make the disclosure unmissable *to anyone who encounters the artifact through the repo*. It does nothing for anyone who encounters the artifact after the Context axis scenario above - clipped, re-encoded, and reposted with the disclosure stripped. Disclosure works well for a

good-faith creator who wants to be transparent and controls distribution; it does close to nothing against a bad-faith actor who never intended to disclose in the first place, or against a good-faith artifact that escapes its original context.

Provenance and content credentials (C2PA, cryptographic signing). I did not verify, and my process log does not show, whether ElevenLabs or D-ID embed C2PA-style content credentials in their default exports. That's itself a finding: if a task specifically built to explore synthetic-media construction and evaluation didn't surface durable provenance metadata as a visible feature of consumer-tier tools, it's reasonable to assume most casual users never encounter it either. Even where provenance metadata exists, it's brittle by design of the distribution pipeline - social platforms routinely strip EXIF and similar metadata on upload for unrelated reasons (file size, privacy), so a credential that isn't independently re-verified by the platform is functionally invisible to end users regardless of whether the original file carried it.

Detection. This is the mitigation I have direct, adversarial evidence against. A purpose-built detector scored my fully synthetic ElevenLabs audio 95% "Likely Authentic," and explained that verdict with language that included a claim about what a spectrogram "would likely reveal" without evidence the tool had actually generated or inspected one. Detection tools are being asked to win a race against generation tools that are, at minimum, keeping pace, and on this one data point, losing outright. A detector that is wrong with high confidence is arguably worse than no detector, because it manufactures false reassurance.

Legal and regulatory regimes. In broad terms: disclosure-mandate laws (requiring synthetic political ads to be labeled), election-specific restrictions on synthetic media within some window of a vote, and non-consensual intimate imagery statutes that criminalize a specific downstream misuse of the technology. All of these regulate *distribution or specific use categories*, not the underlying generation capability - none of them would have stopped me, or anyone, from generating either of my artifacts, because generation itself (with consent, of non-real or licensed content) is not the harm being targeted. That's appropriate for the tools I used, but it also means the law's leverage point (publication, distribution, specific malicious use) arrives after the artifact already exists and, per the Scale axis, potentially after it has already spread.

Platform policy. Major platforms have published synthetic-media policies (labeling requirements, removal of undisclosed political deepfakes, etc.), but enforcement depends on detection working at the moment of upload or report - the same detection step that failed on my artifact. A platform policy that says "we will label or remove synthetic media" is only as strong as its ability to *identify* synthetic media at scale, which is precisely the capability my Task 6 evaluation found to be unreliable even on a clean, single-speaker, no-background-noise sample.

Professional and organizational norms. Journalism, in general, has converged on a norm of not publishing synthetic recreations of real events as if they were documentary footage without clear labeling; advertising and entertainment have more permissive, disclosure-based norms for things like de-aging or posthumous performances. These norms are meaningful within institutions that have reputations to protect and editorial review processes, but they provide no floor at all for individuals or organizations outside those institutions.

Common thread: every mitigation above assumes either a well-intentioned actor who wants to comply, or a detection/verification step that this task's own evidence shows is not reliable. None of them is a solution; each is a partial constraint that raises the cost of misuse without coming close to eliminating it.