

Synthetic Media Policy

Office of University Communications & Marketing

Effective: 09/16/2026 | **Owner:** Director of Communications | **Applies to:** all staff, student workers, and contracted vendors producing content on behalf of the University, including Athletics, Alumni Relations, Admissions, and academic department communications teams.

1. Why this policy exists

This office produces voiceovers, promotional video, social content, and recorded messaging on a routine basis. AI voice and video tools can now generate broadcast-quality narration and lip-synced avatar video in minutes, using stock voices/avatars or, with consent, real staff and students' likenesses. That capability is useful and it is also easy to misuse, sometimes by accident. This policy says what this office will and will not do with it.

2. Permitted uses

The following are permitted **without special review**, provided the disclosure and provenance requirements in Sections 5–6 are met:

- Narrating factual, previously-approved written content (e.g., converting an approved press release, program description, or coach-provided talking points into an AI voiceover) using a **stock voice/avatar**, not a real person's cloned likeness.
- Draft/internal versions of video or audio used for storyboarding, script timing, or A/B testing, never distributed externally.
- Accessibility narration (e.g., generating an audio version of a written article) using a stock voice, clearly labeled as AI-narrated.
- Multilingual narration of already-approved, fact-checked English-language content, using a stock voice in the target language, reviewed by a fluent speaker before publication.

The following are permitted **with review** (Section 7):

- Any use of a **real staff, coach, student-athlete, or student's actual cloned voice or likeness**, even with their consent, in any externally-published material.
- Any synthetic narration of content that has not already gone through this office's standard fact-check process (e.g., quotes, statistics, statements attributed to a named individual).
- Any synthetic media intended to run as paid advertising or appear in fundraising solicitations.

3. Prohibited uses

These are refused outright, regardless of intent, consent, or seniority of the requester:

- Generating synthetic speech or video of **any real, identifiable person** - student, staff, alumnus, donor, or public figure - saying something they did not actually say or verify, even in a draft, even if intended to be corrected before publication. ("We'll just fix the wording later" is not an exception; the artifact should not exist with false content attributed to a real person at any stage.)

- Cloning the voice or likeness of any **minor** (including prospective students under 18 in admissions contexts) under any circumstances, regardless of consent from a parent or guardian.
- Using synthetic media to simulate a **crisis, emergency, or safety announcement** in any testing, demo, or training context where it could plausibly be mistaken for a real alert.
- Producing synthetic content designed to be indistinguishable from an unstaged recording (e.g., faking a “candid” interview or testimonial) even using consenting, disclosed participants - format authenticity matters independently of consent and truthfulness.
- Using a real person’s cloned voice or likeness in any external-facing content **after that person has left the University or withdrawn consent**, regardless of when the original recording or consent was obtained.

4. Consent workflow

When a real person’s voice, likeness, or identifiable characteristics are to be used (cloned, not just appearing on camera in original footage):

1. **Written consent** is obtained by the requesting staff member, using the standard Synthetic Media Consent Form, before any generation begins - not after a draft is made “to see how it sounds.”
2. The form specifies: what will be generated, in what content, for what distribution channels, and for how long the consent is valid (default: one year, renewable).
3. Consent is **revocable at any time**, in writing, effective immediately upon receipt; the requesting office is responsible for removing or ceasing distribution of any content generated from a revoked consent within 5 business days.
4. Signed consent forms are filed with the Director of Communications and retained for the duration of the content’s use plus one year.
5. Consent covers the **specific use described**, not blanket future use - a new campaign requires a new consent instance, even for the same person.

5. Disclosure standards

- Every synthetic voice or video artifact carries an on-screen or spoken disclosure at first appearance (“This video/audio includes AI-generated narration”) - not buried in a credits roll or a linked page.
- File names and any content management system metadata include an AI-GENERATED tag.
- For content likely to be clipped, excerpted, or re-shared (social video especially), the disclosure is burned into the visible frame for at least the first 3 seconds, not relying solely on caption text or a description field that may not travel with re-shares.
- Disclosure is never satisfied by a general disclaimer on a website (“some content on this site may be AI-generated”) applied after the fact to specific artifacts.

6. Provenance requirements

- Where the generation tool supports C2PA content credentials or equivalent embedded metadata, it is enabled by default; staff are not required to manually verify this survives every re-encoding step, but the original master file with credentials intact is retained in the office’s archive as the reference copy.

- The office maintains a log (tool, version, voice/avatar used, date, requester, consent form reference where applicable) for every synthetic media artifact produced, retained for 3 years.

7. Review and approval

- Any “permitted with review” use (Section 2) requires sign-off from the Director of Communications or their designee before generation begins, not before publication - the review gates production, not just release.
- The reviewer checks: is the underlying content fact-checked, is consent on file if a real likeness is involved, does the intended use fall outside any Section 3 prohibition, and is a disclosure plan specified.
- Review turnaround target: 2 business days. Urgent requests may be escalated to the Director directly but are still logged.

8. Incident response

If synthetic media produced under this policy is misused, misattributed, taken out of context, or reaches an unintended audience:

1. The staff member who becomes aware notifies the Director of Communications within 24 hours.
2. The office issues a public correction identifying the specific artifact and clarifying its synthetic nature and original context, posted to the same channel(s) where the misrepresentation occurred, where the office has posting access.
3. If a real person’s likeness was involved without their consent (e.g., a stripped-disclosure re-share), the office notifies that person directly and offers to assist with any platform takedown request they wish to file.
4. The incident and response are logged; recurring incident types trigger a policy review at the next scheduled cycle (Section 9).

9. Policy review

This policy is reviewed annually, or immediately following any incident under Section 8 that reveals a gap the policy did not anticipate.

Limitations

This policy cannot guarantee that synthetic media produced under it will never be misused once it leaves this office’s direct control. Specifically:

- **Disclosure is not durable.** As shown directly in the Task 6 evaluation this policy is based on, a disclosure banner or file-naming convention does not survive being clipped, re-encoded, or reposted. A well-intentioned piece of content that fully complies with Section 5 today can circulate without any trace of that disclosure tomorrow, through no failure of this office’s process. This policy reduces the office’s own liability and sets an internal standard; it does not control what happens to content after it leaves official channels.
- **Detection is not a backstop.** This policy deliberately does not rely on “we can always run it through a detector to check” as a safeguard, because a real detection tool scored genuinely synthetic audio as 95% likely authentic during the research

this policy is grounded in. Staff should not treat a clean detector result - theirs or anyone else's - as confirmation that a piece of media is or isn't synthetic.

- **A well-intentioned employee could still get this wrong.** Someone under deadline pressure might generate a "quick draft" using a real coach's cloned voice before consent paperwork is filed, intending to get consent before publication, and then publish the draft because it was already made and looked fine. The policy names this exact failure mode in Section 3 but cannot physically prevent it; it depends on staff internalizing that the prohibition applies to production, not just release.
- **This policy cannot stop a bad-faith outside actor.** Nothing here constrains someone outside the University who clones a coach's or administrator's voice independently, using publicly available recordings, exactly as described in the Consent-axis scenario this policy is built against. This policy governs what this office produces; it has no jurisdiction over external fabrication, and pairing it with rapid-response incident protocols (Section 8) is a mitigation for that gap, not a solution to it.
- **Enforcement depends on staff self-reporting.** There is no automated system checking whether a given piece of published content followed this consent and review workflow; the policy relies on staff bringing requests through the process rather than generating content independently and treating disclosure as a formality applied afterward. This is an honest reliance on good faith, not a fatal flaw unique to this policy - it is close to inescapable for any policy governing creative production rather than a fully automated pipeline - but it means this policy's real strength depends on training and culture at least as much as on the document itself.